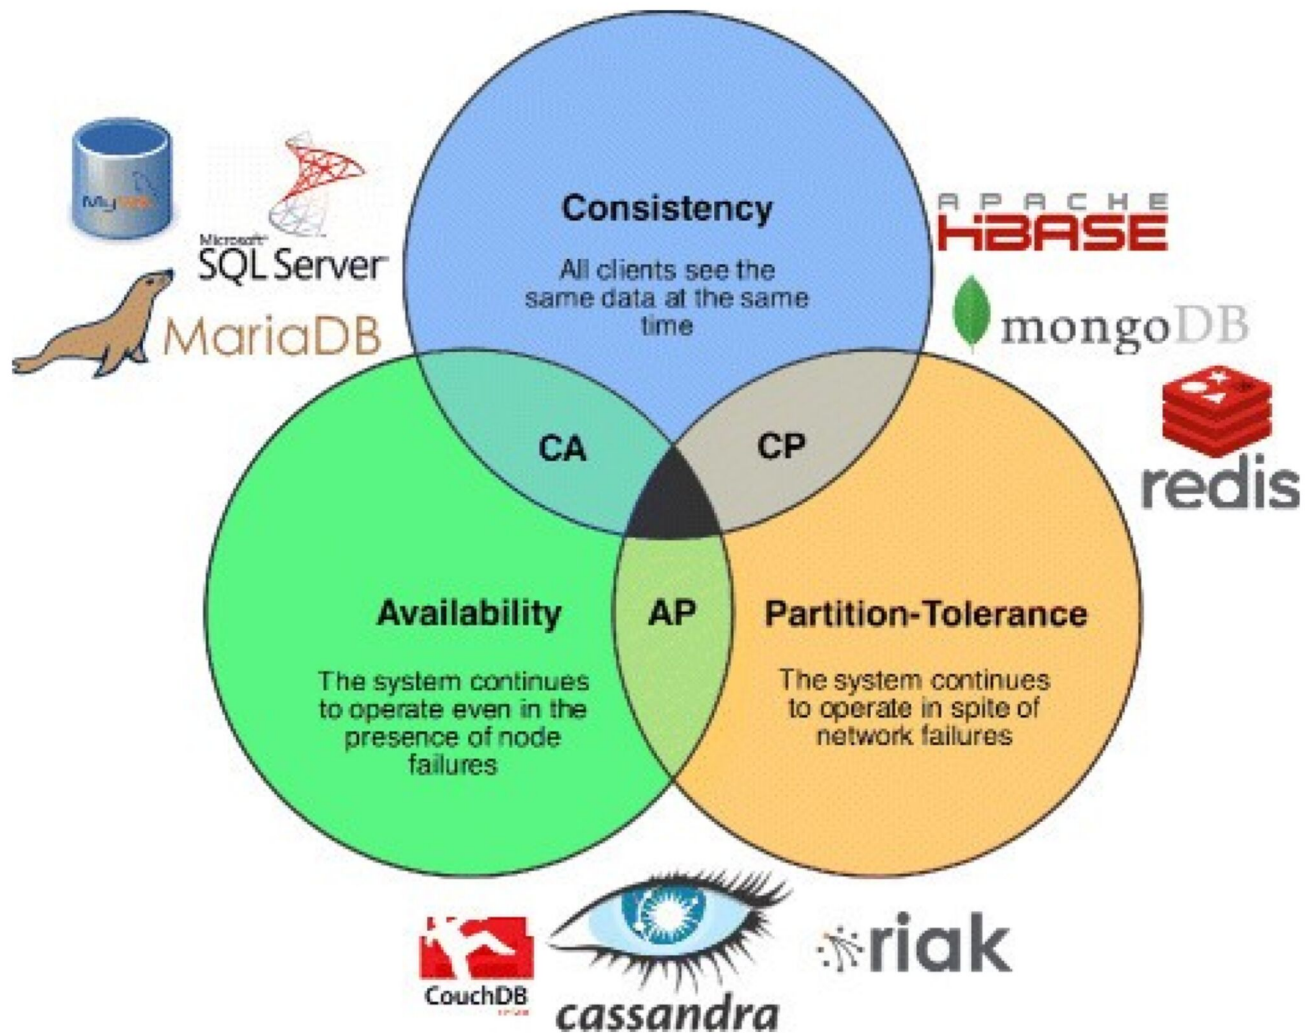


# CAP THEOREM

## (Brewer's Theorem)



CAP — Consistency, Availability, Partition Tolerance

- CAP theorem states that it is impossible for a distributed system to offer more than two out of three following properties:

Simply put,

- Consistency : Every read should receive the recent write or an error.
- Availability : Every request will receive(non-error) response, without the gurantee that the response contains the most recent write.
- Partition Tolerance : The system(cluster) should continue to operate despite any number of communication breakdowns or network failures between nodes. (i.e, a lost connection or temporarily delayed connections between nodes).



Network failures are inevitable in distributed systems , so the system should either choose one of the following in case of network partition :

→ Choosing consistency at the cost of availability :

In this case, an error is returned (say data out-of-date error or data unavailable error) if a particular information is not guaranteed to be up-to-date.

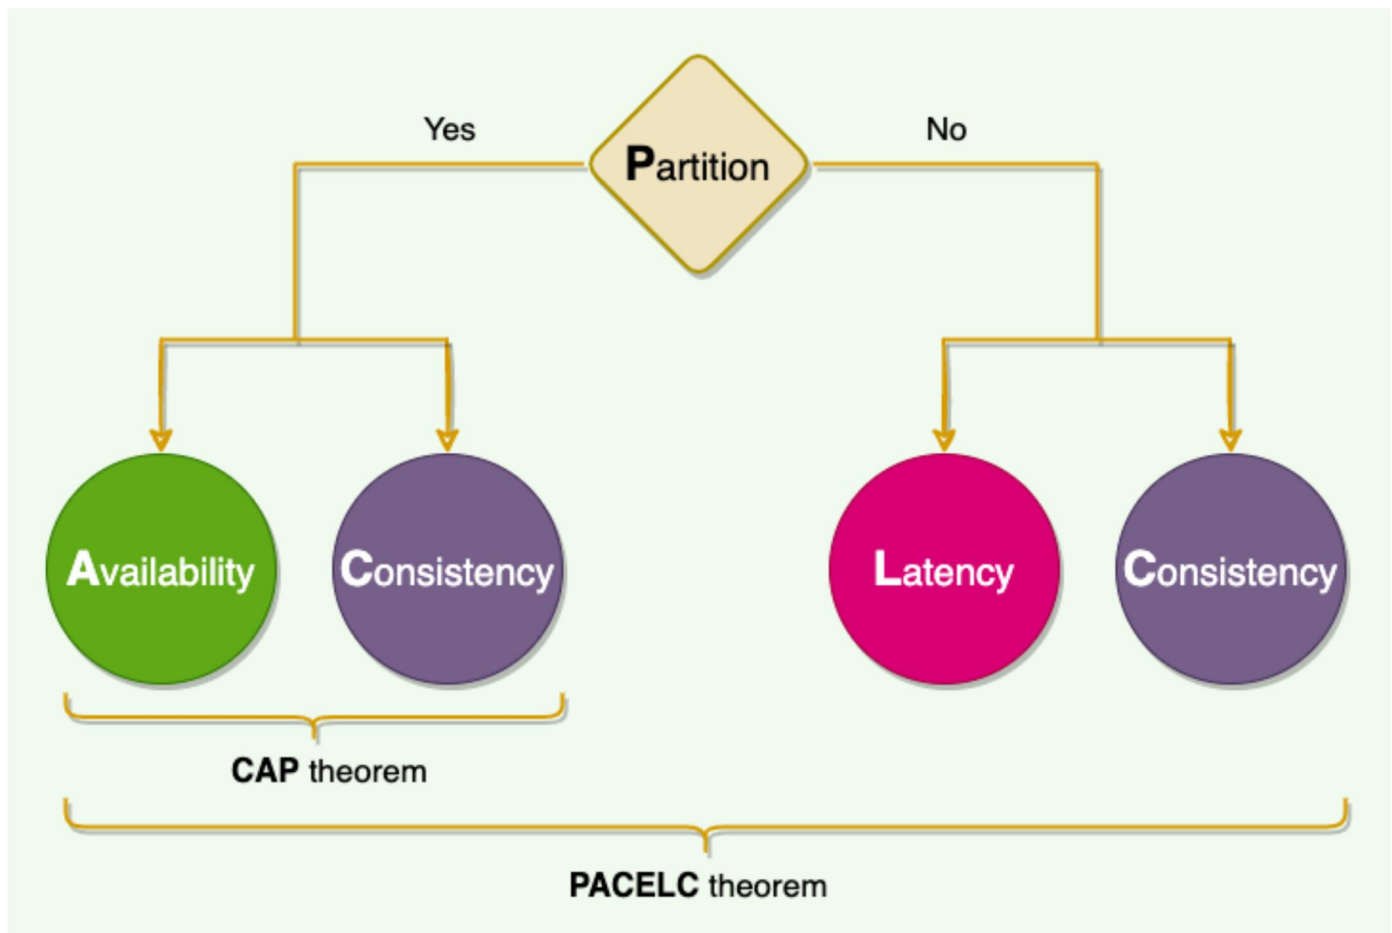
→ Choosing availability at the cost of consistency :

In this case, the system will manage to return the most recent version of data even though it is not guaranteed to be the up-to-date.

NOTE : Database systems designed with traditional ACID guarantees such as RDBMS choose consistency over availability, whereas systems designed around the BASE guarantees, such as NoSQL for example, choose availability over consistency.

# PACELC THEOREM

PACELEC theorem is an extension of CAP theorem.



→ PACELC theorem states that in a distributed system :

In case of network partition(P) the system trades off between availability(A) and consistency(C) Else(E) the system has to tradeoff between latency(L) and consistency(C).

- CAP theorem does not address how a system behaves when there is no partition.  
  
i.e, Let's say the system is serving a request but it's taking 10 mins to return a response. This latency is unacceptable.
- PACELC addresses this key limitation of CAP theorem.
- It states that when the system is functioning normally (without any partition) it either have to choose :
- To stay consistent at all times by ensuring data is updated successfully (which obviously increases latency).
- To return the available data at the moment thereby improving the latency without the guarantee that the data has the most recent updates (non-consistent data).





Santosh Kumar Mishra  
SDE @Microsoft

Follow for more



iamsantoshmishra



iamsantoshmishra



Interview Cafe



Interview Cafe Notes